

In [640...

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

Load MinWage variable

In [641...

```
data = pd.read_csv("./minimum wage.csv")
data['Effective Date'] = pd.to_datetime(data['Effective Date'], format='%m/%d/%Y')
data = data.sort_values(by='Effective Date')

data['Minimum Wage'] = data['Minimum Wage'].replace('[\$,]', '', regex=True)

monthly_range = pd.date_range(start='2003-01-01', end='2024-09-01', freq='MS')
monthly_df = pd.DataFrame({'date': monthly_range})

data = data.set_index('Effective Date')
monthly_df = monthly_df.merge(data[['Minimum Wage']], how='left', left_on='date', right_index=True)

monthly_df['MinWage'] = monthly_df['Minimum Wage'].ffill().bfill()

monthly_df.loc[monthly_df['date'] < '2011-05-01', 'MinWage'] = 8.00

monthly_df = monthly_df.drop(columns=['Minimum Wage'])

monthly_df.head(1000)
```

Out [641...

	date	MinWage
0	2003-01-01	8.00
1	2003-02-01	8.00
2	2003-03-01	8.00
3	2003-04-01	8.00
4	2003-05-01	8.00
...
256	2024-05-01	16.75
257	2024-06-01	17.40
258	2024-07-01	17.40
259	2024-08-01	17.40
260	2024-09-01	17.40

261 rows × 2 columns

load crime_number variable

In [642...

```

crime_data = pd.read_csv("./crimedata_csv_AllNeighbourhoods_AllYears/crime
crime_data['date'] = pd.to_datetime(crime_data[['YEAR', 'MONTH']].assign(
monthly_crime_counts = crime_data.groupby('date').size().reset_index(name=
monthly_df = monthly_df.merge(monthly_crime_counts, how='left', on='date'
monthly_df['crime_number'] = monthly_df['crime_number'].fillna(0)
monthly_df.head(1000)

```

Out [642...

	date	MinWage	crime_number
0	2003-01-01	8.00	4926
1	2003-02-01	8.00	4148
2	2003-03-01	8.00	4550
3	2003-04-01	8.00	4759
4	2003-05-01	8.00	5297
...
256	2024-05-01	16.75	3007
257	2024-06-01	17.40	2810
258	2024-07-01	17.40	3053
259	2024-08-01	17.40	2977
260	2024-09-01	17.40	2566

261 rows × 3 columns

In [643...

```

cpi_data = pd.read_csv("./1810000601-eng (1).csv", header=0, index_col=0)
cpi_data = cpi_data.T
cpi_data = cpi_data.rename(columns={"All-items 8": "CPI"})
cpi_data = cpi_data.reset_index().rename(columns={"index": "date"})
cpi_data['date'] = pd.to_datetime(cpi_data['date'], format='%b-%y')
cpi_data = cpi_data[['date', 'CPI']]
cpi_data.index.name = "index"
monthly_df = monthly_df.merge(cpi_data, how='left', on='date')
monthly_df.head(1000)

```

Out [643...

	date	MinWage	crime_number	CPI
0	2003-01-01	8.00	4926	102.4
1	2003-02-01	8.00	4148	102.0

1	2003-02-01	8.00	4140	102.5
2	2003-03-01	8.00	4550	103.0
3	2003-04-01	8.00	4759	102.3
4	2003-05-01	8.00	5297	102.1
...
256	2024-05-01	16.75	3007	160.6
257	2024-06-01	17.40	2810	160.8
258	2024-07-01	17.40	3053	161.2
259	2024-08-01	17.40	2977	161.3
260	2024-09-01	17.40	2566	161.3

261 rows x 4 columns

In [644...


```
fig, ax1 = plt.subplots(figsize=(12, 6))


color = 'tab:blue'
ax1.set_xlabel('Date')
ax1.set_ylabel('Minimum Wage', color=color)
ax1.plot(monthly_df['date'], monthly_df['MinWage'], color=color)
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx()
color = 'tab:red'
ax2.set_ylabel('Crime Number', color=color)
ax2.plot(monthly_df['date'], monthly_df['crime_number'], color=color)

ax3 = ax1.twinx()
color = 'tab:green'
ax3.spines['right'].set_position(('outward', 60))
ax3.set_ylabel('CPI', color=color)
ax3.plot(monthly_df['date'], monthly_df['CPI'], color=color)

fig.tight_layout()
plt.title('Minimum Wage, Crime Number, and CPI Over Time')
plt.show()
```






Load temperature variable

```
In [645... temperature_data = pd.read_csv("./temperature.csv", header=None, names=["date", "temperature"])
temperature_data['date'] = pd.to_datetime(temperature_data['date'], format='%Y-%m-%d')
merged_df = pd.merge(monthly_df, temperature_data, how='left', on='date')
```

```
In [646... plt.figure(figsize=(12, 6))
plt.plot(merged_df['date'], merged_df['temperature'], color='tab:orange')
plt.title('Temperature Over Time')
plt.xlabel('Date')
plt.ylabel('Temperature')
plt.show()
```



Loaded police officer ratio and weighted clearance rate

```
In [647... police_data = pd.read_csv("./unemployment_police.csv")
police_data['Date'] = pd.to_datetime(police_data['Date'])
police_data.rename(columns={'Date': 'date'}, inplace=True)
merged_df = pd.merge(merged_df, police_data, how='left', on='date')
merged_df.head(6)
```

Out [647...

	date	MinWage	crime_number	CPI	temperature	Police officers per 100,000 population	Weighted clearance rate	Unem
0	2003-01-01	8.0	4926	102.4	7.2	204.2	22.74	
1	2003-02-01	8.0	4148	102.9	5.5	204.2	22.74	
2	2003-03-01	8.0	4550	103.0	7.9	204.2	22.74	
3	2003-04-01	8.0	4759	102.3	9.6	204.2	22.74	
4	2003-05-01	8.0	5297	102.1	15.0	204.2	22.74	
5	2003-06-01	8.0	5199	102.3	17.2	204.2	22.74	

In [648...

```
# print the average of crime number, minimum wage, CPI, temperature, unemp
print("Average of crime number: ", merged_df['crime_number'].mean())
print("Average of minimum wage: ", merged_df['MinWage'].mean())
print("Average of CPI: ", merged_df['CPI'].mean())
print("Average of temperature: ", merged_df['temperature'].mean())
```

Average of crime number: 3434.015325670498
 Average of minimum wage: 10.696360153256707
 Average of CPI: 125.5800766283525
 Average of temperature: 11.418773946360155

In [649...

```
import pandas as pd

def generate_data_description(df):
    stats = df.describe(include='all').transpose()

    stats['unique_values'] = df.nunique()

    stats = stats[['count', 'unique_values', 'mean', 'std', 'min', '25%',

    stats.index.name = 'column_name'

    return stats

data_description = generate_data_description(merged_df)

data_description
```

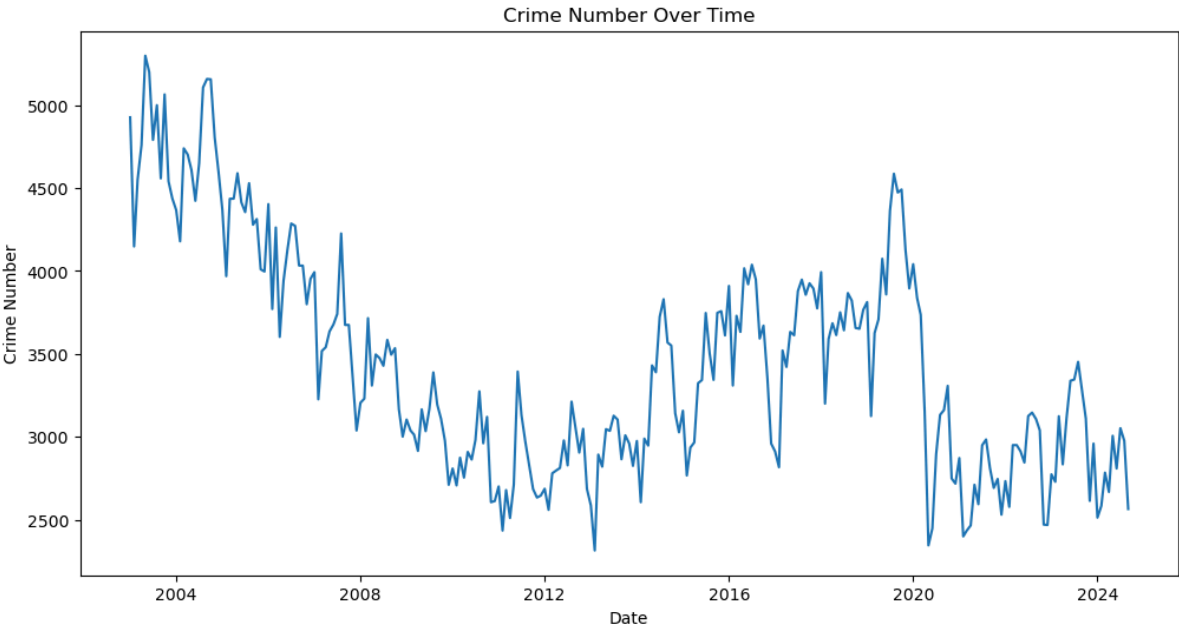
/var/folders/qx/5b_gxq5s5fqc9hdq8fcjn60h0000gn/T/ipykernel_15041/251178911
 1.py:4: FutureWarning: Treating datetime data as categorical rather than nu
 meric in `.describe` is deprecated and will be removed in a future version
 of pandas. Specify `datetime_is_numeric=True` to silence this warning and a
 dopt the future behavior now.
 stats = df.describe(include='all').transpose()

Out [649...

	count	unique_values	mean	std	min	25%	50%
column_name							
date	261	261	NaN	NaN	NaN	NaN	NaN
MinWage	261.0	11	10.69636	2.960168	8.0	8.0	10
crime_number	261.0	245	3434.015326	672.739325	2316.0	2913.0	331
CPI	261.0	206	125.580077	15.707729	102.1	113.7	12
temperature	261.0	136	11.418774	5.351577	2.3	6.5	1
Police officers per 100,000 population	261.0	20	204.136782	15.397368	183.8	191.5	19
Weighted clearance rate	261.0	19	26.333678	2.262443	22.74	24.53	25
Vancouver Unemployment Rate	261.0	56	6.164368	1.572723	3.2	4.8	

In [650...

```
# plot the crime number data over time
plt.figure(figsize=(12, 6))
plt.plot(merged_df['date'], merged_df['crime_number'], color='tab:blue')
plt.title('Crime Number Over Time')
plt.xlabel('Date')
plt.ylabel('Crime Number')
plt.show()
```



In [651...

```
# visualize Police officers per 100,000 population and weighted clearance rate
fig, ax1 = plt.subplots(figsize=(12, 6))

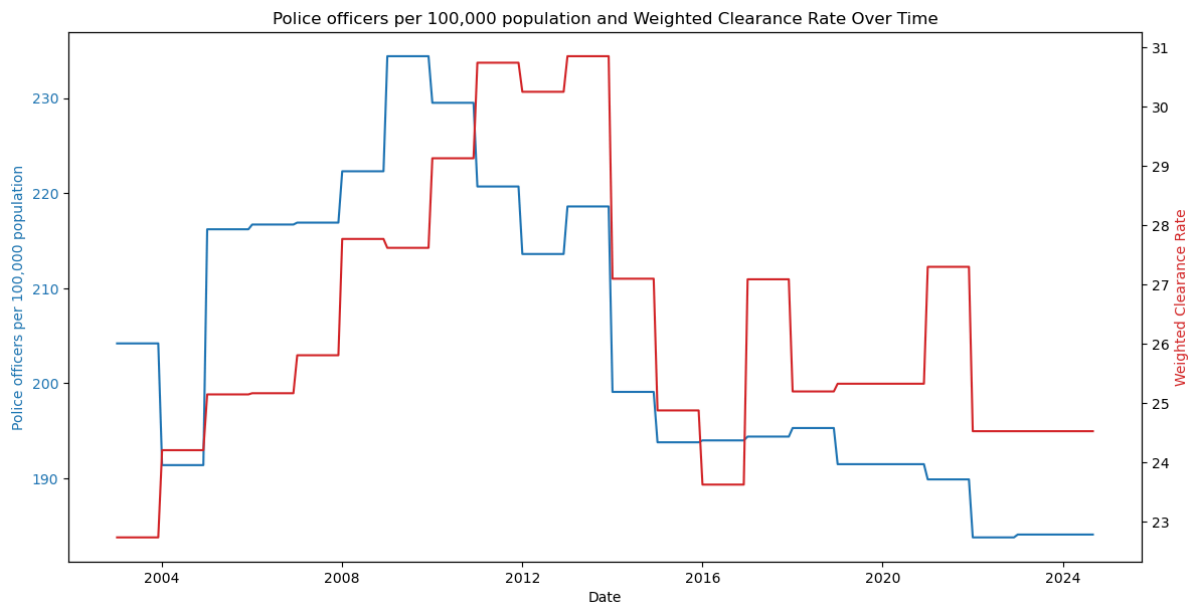
color = 'tab:blue'
ax1.set_xlabel('Date')
ax1.set_ylabel('Police officers per 100,000 population', color=color)
```

```
ax1.plot(merged_df['date'], merged_df['Police officers per 100,000 population'], color='blue')
ax1.tick_params(axis='y', labelcolor='blue')

ax2 = ax1.twinx()
color = 'tab:red'

ax2.set_ylabel('Weighted Clearance Rate', color=color)
ax2.plot(merged_df['date'], merged_df['Weighted clearance rate'], color=color)

fig.tight_layout()
plt.title('Police officers per 100,000 population and Weighted Clearance Rate Over Time')
plt.show()
```



In [652...

```
import matplotlib.pyplot as plt
import seaborn as sns

merged_df['year'] = merged_df['date'].dt.year

fig, ax = plt.subplots(2, 3, figsize=(12, 12))

sns.boxplot(data=merged_df, x='year', y='Vancouver Unemployment Rate', ax=ax[0, 0])
ax[0, 0].set_title('Vancouver Unemployment Rate')
ax[0, 0].set_xlabel('Year')
ax[0, 0].set_ylabel('Unemployment Rate (%)')
ax[0, 0].tick_params(axis='x', rotation=45)

sns.boxplot(data=merged_df, x='year', y='MinWage', ax=ax[0, 1], palette='Reds')
ax[0, 1].set_title('Minimum Wage')
ax[0, 1].set_xlabel('Year')
ax[0, 1].set_ylabel('Minimum Wage (CAD)')
ax[0, 1].tick_params(axis='x', rotation=45)

sns.boxplot(data=merged_df, x='year', y='CPI', ax=ax[1, 0], palette='Reds')
ax[1, 0].set_title('CPI')
ax[1, 0].set_xlabel('Year')
ax[1, 0].set_ylabel('CPI')
ax[1, 0].tick_params(axis='x', rotation=45)

sns.boxplot(data=merged_df, x='year', y='temperature', ax=ax[1, 1], palette='Reds')
ax[1, 1].set_title('Temperature')
```

```

ax[1, 1].set_xlabel('Year')
ax[1, 1].set_ylabel('Temperature (°C)')
ax[1, 1].tick_params(axis='x', rotation=45)

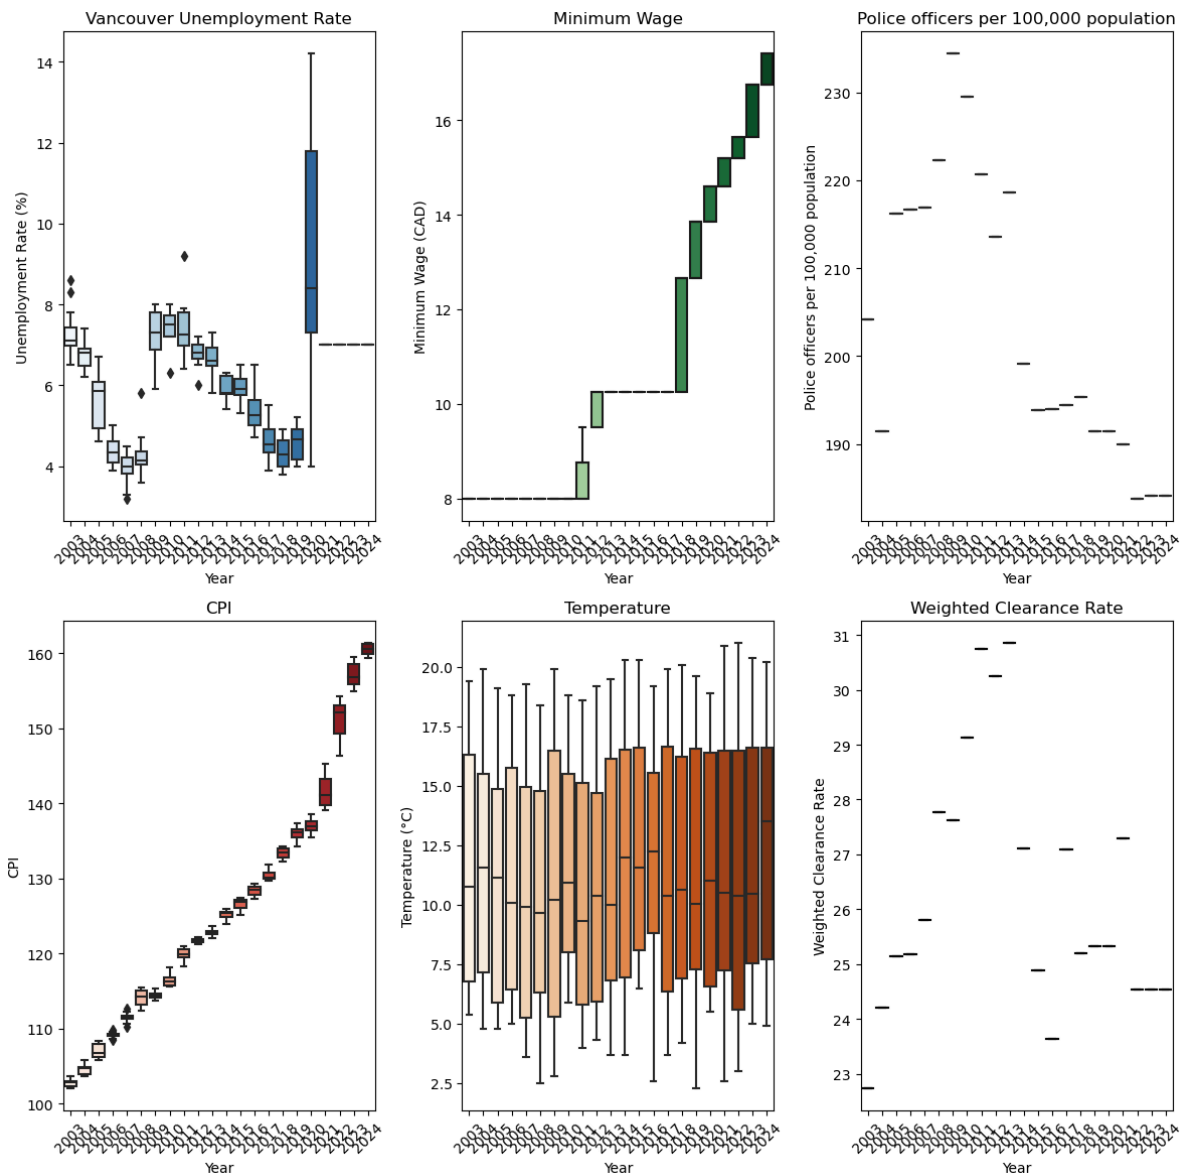
sns.boxplot(data=merged_df, x='year', y='Police officers per 100,000 popu
ax[0, 2].set_title('Police officers per 100,000 population')
ax[0, 2].set_xlabel('Year')
ax[0, 2].set_ylabel('Police officers per 100,000 population')
ax[0, 2].tick_params(axis='x', rotation=45)

sns.boxplot(data=merged_df, x='year', y='Weighted clearance rate', ax=ax[1
ax[1, 2].set_title('Weighted Clearance Rate')
ax[1, 2].set_xlabel('Year')
ax[1, 2].set_ylabel('Weighted Clearance Rate')
ax[1, 2].tick_params(axis='x', rotation=45)

plt.suptitle('')
plt.tight_layout()

# Show the plot
plt.show()

```




```
merged_df_arima = merged_df.copy()
merged_df_xgb = merged_df.copy()
```

Method 1: SARIMA

In [654...

```
from statsmodels.tsa.statespace.sarimax import SARIMAX
from sklearn.metrics import mean_absolute_percentage_error
import matplotlib.pyplot as plt
import pandas as pd

train_data = merged_df_arima[(merged_df_arima['date'] < '2019-01-01')]
val_data = merged_df_arima[(merged_df_arima['date'] >= '2019-01-01') & (merged_df_arima['date'] < '2020-01-01')]

crime_series_train = train_data.set_index('date')['crime_number'].asfreq('MS')
crime_series_val = val_data.set_index('date')['crime_number'].asfreq('MS')

sarima_model = SARIMAX(crime_series_train, order=(1, 0, 0), seasonal_order=(0, 0, 0, 0))

forecast_val = sarima_model.predict(start=crime_series_val.index[0], end=crime_series_val.index[-1])

mape_val = mean_absolute_percentage_error(crime_series_val, forecast_val)
print(f"Validation MAPE: {mape_val * 100:.2f}%")

full_train_data = merged_df_arima[merged_df_arima['date'] < '2020-01-01']
sarima_model_full = SARIMAX(full_train_data, order=(1, 0, 0), seasonal_order=(0, 0, 0, 0))

forecast_start = '2020-01-01'
forecast_end = merged_df_arima['date'].max()
forecast_test = sarima_model_full.predict(start=forecast_start, end=forecast_end)



forecast_test_df = forecast_test.reset_index()
forecast_test_df.columns = ['date', 'predicted_crime_number_arima']
merged_df_arima = pd.merge(merged_df_arima, forecast_test_df, how='left', on='date')

plt.figure(figsize=(14, 7))
plt.plot(merged_df_arima['date'], merged_df_arima['crime_number'], label='Actual Crime Number')
plt.plot(merged_df_arima['date'], merged_df_arima['predicted_crime_number_arima'], label='Predicted Crime Number')
plt.axvline(x=pd.to_datetime('2020-01-01'), color='green', linestyle='--')
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("Actual vs Predicted Crime Numbers (SARIMA)")
plt.legend()
plt.show()

plt.figure(figsize=(14, 7))
plt.plot(crime_series_val.index, crime_series_val, label="Actual Crime Number")
plt.plot(crime_series_val.index, forecast_val, label="Predicted Crime Number")
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("SARIMA Validation Predictions")
plt.legend()
plt.show()
```

Validation MAPE: 9.95%

Actual vs Predicted Crime Numbers (SARIMA)



In [655...

```

from sklearn.metrics import mean_absolute_percentage_error

test_data_actual = merged_df_arima.loc[merged_df_arima['date'] >= '2020-03-01']
test_data_predicted = merged_df_arima.loc[merged_df_arima['date'] >= '2020-03-01']

valid_indices = test_data_actual.notna() & test_data_predicted.notna()
test_data_actual = test_data_actual[valid_indices]
test_data_predicted = test_data_predicted[valid_indices]

mape_test = mean_absolute_percentage_error(test_data_actual, test_data_predicted)
print(f"Test MAPE: {mape_test * 100:.2f}%")

percent_diff = ((test_data_actual - test_data_predicted) / test_data_actual)
print(f"Average percent difference after 2020-03-01: {percent_diff:.2f}%")


plt.figure(figsize=(14, 7))
plt.plot(merged_df_arima.loc[merged_df_arima['date'] >= '2020-03-01', 'actual'], label='Actual')
plt.plot(merged_df_arima.loc[merged_df_arima['date'] >= '2020-03-01', 'predicted'], label='Predicted')
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("Actual vs Predicted Crime Numbers (Test Period - SARIMA)")
plt.legend()

```

```
plt.show()
```

Test MAPE: 25.31%

Average percent difference after 2020-03-01: -25.31%



Method 2: XGBoost

In [656...

```
import xgboost as xgb
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import mean_absolute_percentage_error

def create_lagged_features(df, target_col, lags):
    df = df.copy()
    for lag in range(1, lags + 1):
        df[f'{target_col}_lag_{lag}'] = df[target_col].shift(lag)
    return df

train_data = merged_df_xgb[merged_df_xgb['date'] < '2019-01-01']
val_data = merged_df_xgb[(merged_df_xgb['date'] >= '2019-01-01') & (merged_df_xgb['date'] < '2020-01-01')]

train_data = create_lagged_features(train_data, 'crime_number', 3)
val_data = create_lagged_features(val_data, 'crime_number', 3)

val_data.fillna(method='ffill', inplace=True)
val_data.fillna(val_data.mean(numeric_only=True), inplace=True)

train_data['month'] = train_data['date'].dt.month
train_data['year'] = train_data['date'].dt.year
val_data['month'] = val_data['date'].dt.month
val_data['year'] = val_data['date'].dt.year

features = [
    'crime_number_lag_1', 'crime_number_lag_2', 'crime_number_lag_3',
    'month', 'year', 'Police officers per 100,000 population',
    'Weighted clearance rate', 'Vancouver Unemployment Rate',
    'CPI', 'MinWage'
]
```

```

X_train = train_data[features]
y_train = train_data['crime_number']
X_val = val_data[features]
y_val = val_data['crime_number']

xgb_model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=1000)
xgb_model.fit(X_train, y_train)

val_predictions = xgb_model.predict(X_val)

plt.figure(figsize=(14, 7))

plt.plot(val_data['date'], y_val, label="Actual Crime Number (Validation)")
plt.plot(val_data['date'], val_predictions, label="Predicted Crime Number")

plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("XGBoost Validation Predictions")
plt.legend()
plt.show()

mape_val = mean_absolute_percentage_error(y_val, val_predictions)
print(f"Validation MAPE: {mape_val * 100:.2f}%")

full_train_data = merged_df_xgb[merged_df_xgb['date'] < '2020-01-01']
full_train_data = create_lagged_features(full_train_data, 'crime_number',
full_train_data['month'] = full_train_data['date'].dt.month
full_train_data['year'] = full_train_data['date'].dt.year

X_full_train = full_train_data[features]
y_full_train = full_train_data['crime_number']

xgb_model.fit(X_full_train, y_full_train)

test_data = merged_df_xgb[merged_df_xgb['date'] >= '2020-01-01'].copy()
test_data = create_lagged_features(test_data, 'crime_number', 3).fillna(0)
test_data['month'] = test_data['date'].dt.month
test_data['year'] = test_data['date'].dt.year

test_predictions = []
for i in range(len(test_data)):
    X_test_row = test_data[features].iloc[i].values.reshape(1, -1)
    pred = xgb_model.predict(X_test_row)[0]
    test_predictions.append(pred)

    if i + 1 < len(test_data):
        test_data.loc[test_data.index[i + 1], 'crime_number_lag_1'] = pred
        test_data.loc[test_data.index[i + 1], 'crime_number_lag_2'] = test_predictions[i]
        test_data.loc[test_data.index[i + 1], 'crime_number_lag_3'] = test_predictions[i-1]


test_data['predicted_crime_number_xgb'] = test_predictions

merged_df_xgb = pd.merge(merged_df_xgb, test_data[['date', 'predicted_crime_number_xgb']],
                        on='date', how='left')

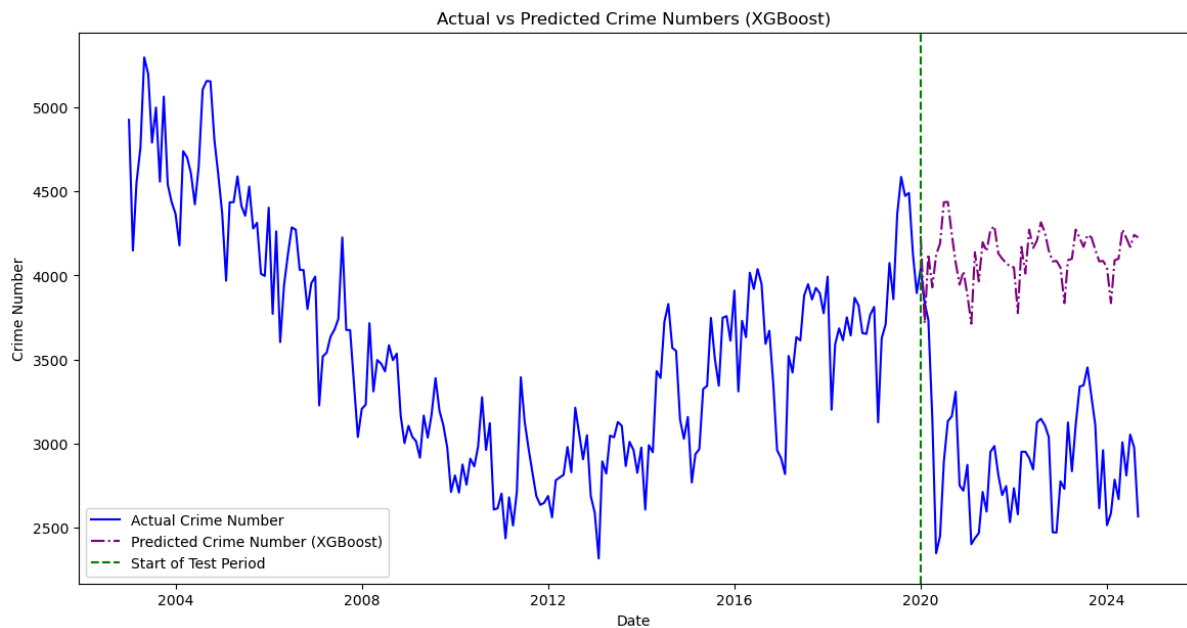
plt.figure(figsize=(14, 7))
plt.plot(merged_df_xgb['date'], merged_df_xgb['crime_number'], label="Actual Crime Number")
plt.plot(merged_df_xgb['date'], merged_df_xgb['predicted_crime_number_xgb'], label="Predicted Crime Number")
plt.axvline(x=pd.to_datetime('2020-01-01'), color='green', linestyle='--')
plt.xlabel("Date")

```

```
plt.ylabel("Crime Number")
plt.title("Actual vs Predicted Crime Numbers (XGBoost)")
plt.legend()
plt.show()
```



Validation MAPE: 8.94%



In [657...

```
from sklearn.metrics import mean_absolute_percentage_error

# Filter test data after 2020-03-01 (real test data period)
test_data_actual = merged_df_xgb.loc[merged_df_xgb['date'] >= '2020-03-01']
test_data_predicted = merged_df_xgb.loc[merged_df_xgb['date'] >= '2020-03-01']

# Drop NaN values
valid_indices = test_data_actual.notna() & test_data_predicted.notna()
test_data_actual = test_data_actual[valid_indices]
test_data_predicted = test_data_predicted[valid_indices]

# Calculate MAPE
```

```

mape_test = mean_absolute_percentage_error(test_data_actual, test_data_predicted)
print(f"Test MAPE: {mape_test * 100:.2f}%")

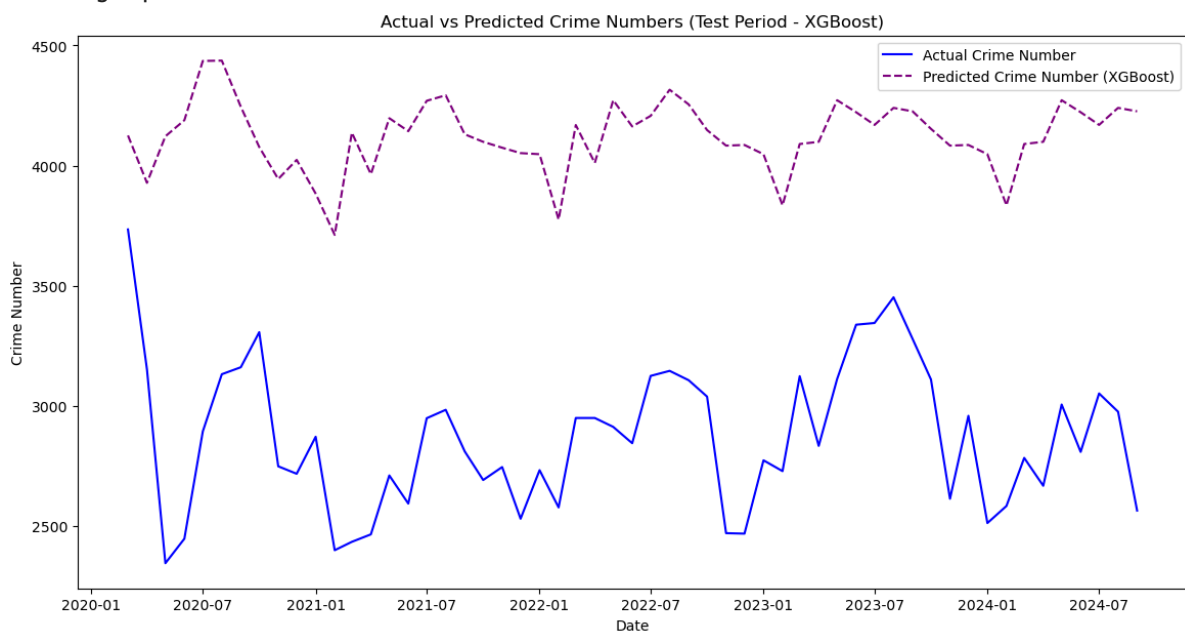
# Calculate percentage difference
percent_diff = ((test_data_actual - test_data_predicted) / test_data_actual) * 100
print(f"Average percent difference after 2020-03-01: {percent_diff:.2f}%")

# Optional: Visualize actual vs predicted values for the test period
plt.figure(figsize=(14, 7))
plt.plot(merged_df_xgb.loc[merged_df_xgb['date'] >= '2020-03-01', 'date'], merged_df_xgb.loc[merged_df_xgb['date'] >= '2020-03-01', 'crime_number'], label="Actual Crime Number")
plt.plot(merged_df_xgb.loc[merged_df_xgb['date'] >= '2020-03-01', 'date'], merged_df_xgb.loc[merged_df_xgb['date'] >= '2020-03-01', 'predicted_crime_number_xgb'], label="Predicted Crime Number (XGBoost)")
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("Actual vs Predicted Crime Numbers (Test Period - XGBoost)")
plt.legend()
plt.show()

```

Test MAPE: 44.95%

Average percent difference after 2020-03-01: -44.95%



In [658...

```

# put Actual vs Predicted Crime Numbers (XGBoost), XGBoost Validation Predictions
# 1 figure with 4 subplots
plt.figure(figsize=(14, 14))

# Actual vs Predicted Crime Numbers (XGBoost)
plt.subplot(2, 2, 1)

plt.plot(merged_df_xgb['date'], merged_df_xgb['crime_number'], label="Actual Crime Number")
plt.plot(merged_df_xgb['date'], merged_df_xgb['predicted_crime_number_xgb'], label="Predicted Crime Number (XGBoost)")
plt.axvline(x=pd.to_datetime('2020-01-01'), color='green', linestyle='--')
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("Actual vs Predicted Crime Numbers (XGBoost)")
plt.legend()

# XGBoost Validation Predictions
plt.subplot(2, 2, 2)

plt.plot(val_data['date'], y_val, label="Actual Crime Number (Validation)")
plt.plot(val_data['date'], val_predictions, label="Predicted Crime Number (XGBoost Validation)")

```

```

plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("XGBoost Validation Predictions")
plt.legend()

# Actual vs Predicted Crime Numbers (SARIMA)
plt.subplot(2, 2, 3)

plt.plot(merged_df_arima['date'], merged_df_arima['crime_number'], label='Actual Crime Number (Validation)')
plt.plot(merged_df_arima['date'], merged_df_arima['predicted_crime_number'], label='Predicted Crime Number (Validation, XGBoost)')
plt.axvline(x=pd.to_datetime('2020-01-01'), color='green', linestyle='--')
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("Actual vs Predicted Crime Numbers (SARIMA)")
plt.legend()

# SARIMA Validation Predictions
plt.subplot(2, 2, 4)

plt.plot(val_data['date'], y_val, label="Actual Crime Number (Validation)")
plt.plot(val_data['date'], forecast_val, label="Predicted Crime Number (Validation, SARIMA)")
plt.xlabel("Date")
plt.ylabel("Crime Number")
plt.title("SARIMA Validation Predictions")
plt.legend()

plt.tight_layout()
plt.show()

```

